

Milestone 8 Bewijs in PDF

Sam Rotthier - TIFL

Overzicht vergelijking:

Tabel info voor partitionering:

Query:

```
select segment_name, segment_type, sum(bytes/1024/1024) MB ,  
(select COUNT(*) FROM monster) as table_count  
from dba_segments  
where segment_name= 'MONSTER'  
group by segment_name, segment_type;
```

	SEGMENT_NAME	SEGMENT_TYPE	MB	TABLE_COUNT
1	MONSTER	TABLE	17.0625	400000

Query:

```
SELECT p.playerid, m.monstername, ROUND(AVG(m."level")) AS "Average_time_played"  
FROM player p  
      JOIN team t ON p.playerid = t.PLAYER_PLAYERID  
      JOIN monster m ON m.team_teamid = t.teamid  
WHERE t.teamid < 25  
GROUP BY p.playerid, m.monstername  
ORDER BY p.playerid;
```

Explain plan

<pre>SELECT p.playerid, m.monstername, ROUND(AVG(m."level")) AS "Average_time_played" FROM player p JOIN team t ON p.playerid = t.PLAYER_PLAYERID JOIN monster m ON m.team_teamid = t.teamid WHERE t.teamid < 25 GROUP BY p.playerid, m.monstername ORDER BY p.playerid;</pre>					
Operation	Params	Rows	Total Cost	Raw Desc	
← Select		14	588.0	cpu_cost = 154105478, io_...	
Group By (SORT GROUP BY)		14	588.0	cpu_cost = 154105478, io_...	
Hash Join		2310	587.0	cpu_cost = 123335724, io_...	
Access (TABLE ACCESS BY table: TEAM;		24	3.0	cpu_cost = 32095, io_cost = 3	
Index Scan (INDEX RAN index: TEAM_PK;		24	2.0	cpu_cost = 19243, io_cost = 2	
Full Scan (TABLE ACCESS F table: MONSTER;		7113	584.0	cpu_cost = 121988729, io_...	

NA partitioning:

Partitie script + uitleg partitie sleutel

```
DROP TABLE monster CASCADE CONSTRAINTS PURGE;
CREATE TABLE monster (
    monsterid INTEGER GENERATED ALWAYS AS IDENTITY
    CONSTRAINT software_pk PRIMARY KEY,
    monstername VARCHAR2(500) NOT NULL,
    health INTEGER NOT NULL,
    "level" INTEGER,
    canevoive CHAR(1),
    team_teamid INTEGER NOT NULL
)
PARTITION BY RANGE (team_teamid)
INTERVAL (50)
(
    partition team_50 VALUES LESS THAN (50)
);
ALTER TABLE monster
ADD CONSTRAINT monster_team_fk FOREIGN KEY (team_teamid)
REFERENCES team ( teamid );
```

Tabel info NA partitioning:

Query:

```
select segment_name,segment_type, sum(bytes/1024/1024) MB ,
(select COUNT(*) FROM monster) as table_count
from dba_segments
where segment_name= 'MONSTER'
group by segment_name,segment_type;
```

	SEGMENT_NAME	SEGMENT_TYPE	MB	TABLE_COUNT
1	MONSTER	TABLE	0.0625	400000
2	MONSTER	TABLE PARTITION	72	400000

Explain plan na partitionering

```
SELECT p.playerid, m.monstername, ROUND(AVG(m."level")) AS "Average_time_played"
FROM player p
      JOIN team t ON p.playerid = t.PLAYER_PLAYERID
      JOIN monster m ON m.team_teamid = t.teamid
WHERE t.teamid < 25
GROUP BY p.playerid, m.monstername
ORDER BY p.playerid;
```

Operation	Params	Rows	Total Cost	Raw Desc
↩ Select		14	280.0	cpu_cost = 89698435, io_c...
↳ Group By (SORT GROUP BY)		14	280.0	cpu_cost = 89698435, io_c...
↳ Hash Join		3361	279.0	cpu_cost = 58317659, io_c...
↳ Unknown (PART JOIN FILTER CRE		24	3.0	cpu_cost = 32095, io_cost ...
↳ Access (TABLE ACCESS BY table: TEAM;		24	3.0	cpu_cost = 32095, io_cost ...
↳ Index Scan (INDEX RAI index: TEAM_PK;		24	2.0	cpu_cost = 19243, io_cost ...
↳ Unknown (PARTITION RANGE SIN		3497	276.0	cpu_cost = 57332264, io_c...
↳ Full Scan (TABLE ACCESS table: MONSTER;		3497	276.0	cpu_cost = 57332264, io_c...

Blooper:

Voor:

```
SELECT p.playerid, p.name, ROUND(AVG(m."level")) AS "Average_monster_level"
FROM player p
      JOIN team t ON p.playerid = t.PLAYER_PLAYERID
      JOIN monster m ON m.team_teamid = t.teamid
WHERE m."level" > 25
GROUP BY p.playerid, p.name;
```

Operation	Params	Rows	Total Cost	Raw Desc
Select		20	599.0	cpu_cost = 426815600, i...
Merge Join		20	599.0	cpu_cost = 426815600, i...
Index Scan (TABLE ACCESS BY INDEX)	table: PLAYER;	20	2.0	cpu_cost = 22643, io_cost...
Full Index Scan (INDEX FULL SCAN)	index: PLAYER_PK;	20	1.0	cpu_cost = 11121, io_cost...
Sort (SORT JOIN)		20	597.0	cpu_cost = 426792957, i...
Access (VIEW)		20	596.0	cpu_cost = 397182221, i...
Group By (HASH GROUP BY)		20	596.0	cpu_cost = 397182221, i...
Hash Join		269799	588.0	cpu_cost = 148268422, i...
Full Scan (TABLE ACCESS)	table: TEAM;	800	3.0	cpu_cost = 203607, io_co...
Full Scan (TABLE ACCESS)	table: MONSTER;	269799	584.0	cpu_cost = 120364915, i...

Na:

```
SELECT p.playerid, p.name, ROUND(AVG(m."level")) AS "Average_monster_level"
FROM player p
      JOIN team t ON p.playerid = t.PLAYER_PLAYERID
      JOIN monster m ON m.team_teamid = t.teamid
WHERE m."level" > 25
GROUP BY p.playerid, p.name;
```

Operation	Params	Rows	Total Cost	Raw Desc
Select		20	917.0	cpu_cost = 2755834168,...
Merge Join		20	917.0	cpu_cost = 2755834168,...
Index Scan (TABLE ACCESS BY INDEX)	table: PLAYER;	20	2.0	cpu_cost = 22643, io_cos...
Full Index Scan (INDEX FULL SCAN)	index: PLAYER_PK;	20	1.0	cpu_cost = 11121, io_cos...
Sort (SORT JOIN)		20	915.0	cpu_cost = 2755811526,...
Access (VIEW)		20	914.0	cpu_cost = 2726200789,...
Group By (HASH GROUP BY)		20	914.0	cpu_cost = 2726200789,...
Hash Join		230209	907.0	cpu_cost = 2511842411,...
Full Scan (TABLE ACCESS)	table: TEAM;	800	3.0	cpu_cost = 203607, io_c...
Unknown (PARTITION RANG		230209	903.0	cpu_cost = 2487897904,...
Full Scan (TABLE ACCESS)	table: MONSTER;	230209	903.0	cpu_cost = 2487897904,...

Conclusie:

In de blooper kunnen we zien dat wanneer je een index op een verkeerd attribuut zet (of een verkeerde filter toepast) dat dit tegenovergesteld werkt en de query langer laat duren.

Dit komt doordat het programma nu over meerdere files gaat moeten zoeken in plaats van (bij een goede partitionering) minder.

Door de blooper zie je duidelijk hoe de theorie de voor en nadelen van een partitie bespreekt.

Bij een goede parititie (en filtering) gaat het syteem net over minder files moeten zoeken en hierdoor veel tijd winnen.